Monday 3rd of October 4:30pm

Should we use multiple classifiers and compare them?

Should we try to improve a classifier and see how it affects the results?

Should we use multiple evaluation methods and compare them?

Should we use different methods of creating training and development sets?

Should we be looking at individual tweets?

This classifier is really good at predicting for this class according to these evaluation methods because x.

Don't test on the training set!
Just something slightly outside of the box, look at the discussion forum for examples of creativity. Creative baseline for evaluation, really good bibliography.

Talk about noise. Talk about number of features used.

Feature engineering is **not** reducing the number of features, that is feature selection. Feature engineering is only when you select additional features by some method.